

Handling sequences with PyTorch



INTERMEDIATE DEEP LEARNING WITH PYTORCH



Michał Oleszak
Machine Learning Engineer

Sequential data

- Ordered in time or space
- Order of the data points contains dependencies between them
- Examples of sequential data:
 - Time series
 - Text
 - Audio waves



Electricity consumption prediction


- Task: predict future electricity consumption based on past patterns
- Electricity consumption dataset:

	timestamp	consumption
0	2011-01-01 00:15:00	-0.704319
1	2011-01-01 00:30:00	-0.704319
...
140254	2014-12-31 23:45:00	-0.095751
140255	2015-01-01 00:00:00	-0.095751

¹ Trindade, Artur. (2015). ElectricityLoadDiagrams20112014. UCI Machine Learning Repository.
<https://doi.org/10.24432/C58C86>.


Train-test split

- No random splitting for time series!
- **Look-ahead bias:** model has info about the future
- Solution: split by time



Creating sequences

- Sequence length = number of data points in one training example
 - $24 \times 4 = 96$ -> consider last 24 hours
- Predict single next data point



Creating sequences in Python

```
import numpy as np

def create_sequences(df, seq_length):
    xs, ys = [], []
    for i in range(len(df) - seq_length):
        x = df.iloc[i:(i+seq_length), 1]
        y = df.iloc[i+seq_length, 1]
        xs.append(x)
        ys.append(y)
    return np.array(xs), np.array(ys)
```

- Take data and sequence length as inputs
- Initialize inputs and targets lists
- Iterate over data points
- Define inputs and target
- Append to pre-initialized lists
- Return inputs and targets as NumPy arrays

TensorDataset

Create training examples

```
X_train, y_train = create_sequences(train_data, seq_length)  
print(X_train.shape, y_train.shape)
```

```
(34944, 96) (34944,)
```

Convert them to a Torch Dataset

```
from torch.utils.data import TensorDataset  
  
dataset_train = TensorDataset(  
    torch.from_numpy(X_train).float(),  
    torch.from_numpy(y_train).float(),  
)
```

Applicability to other sequential data

Same techniques are applicable to other sequences:

- Large Language Models
- Speech recognition

Let's practice!

INTERMEDIATE DEEP LEARNING WITH PYTORCH

Recurrent Neural Networks


INTERMEDIATE DEEP LEARNING WITH PYTORCH




Michał Oleszak
Machine Learning Engineer

Recurrent neuron


- Feed-forward networks
- RNNs: have connections pointing back
- Recurrent neuron:
 - Input x
 - Output y
 - Hidden state h
- In PyTorch: `nn.RNN()`




Unrolling recurrent neuron through time




Unrolling recurrent neuron through time



Unrolling recurrent neuron through time




Deep RNNs




Sequence-to-sequence architecture

- Pass sequence as input, use the entire output sequence
- **Example:** Real-time speech recognition




Sequence-to-vector architecture

- Pass sequence as input, use only the last output
- **Example:** Text topic classification




Vector-to-sequence architecture

- Pass single input, use the entire output sequence
- **Example:** Text generation



Encoder-decoder architecture

- Pass entire input sequence, only then start using output sequence
- **Example:** Machine translation



RNN in PyTorch

```
class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.rnn = nn.RNN(
            input_size=1,
            hidden_size=32,
            num_layers=2,
            batch_first=True,
        )
        self.fc = nn.Linear(32, 1)

    def forward(self, x):
        h0 = torch.zeros(2, x.size(0), 32)
        out, _ = self.rnn(x, h0)
        out = self.fc(out[:, -1, :])
        return out
```

- Define model class with `__init__` method
- Define recurrent layer, `self.rnn`
- Define linear layer, `fc`
- In `forward()`, initialize first hidden state to zeros
- Pass input and first hidden state through RNN layer
- Select last RNN's output and pass it through linear layer

Let's practice!

INTERMEDIATE DEEP LEARNING WITH PYTORCH

LSTM and GRU cells

INTERMEDIATE DEEP LEARNING WITH PYTORCH




Michał Oleszak


Machine Learning Engineer

Short-term memory problem

- RNN cells maintain memory via hidden state
- This memory is very short-term
- Two more powerful cells solve the problem:
 - LSTM (Long Short-Term Memory) cell
 - GRU (Gated Recurrent Unit) cell




RNN cell



- Two inputs:
 - current input data x
 - previous hidden state h
- Two outputs:
 - current output y
 - next hidden state h

LSTM cell



- Three inputs and outputs (two hidden states):
 - h : short-term state
 - c : long-term state
- Three "gates":
 - **Forget gate**: what to remove from long-term memory
 - **Input gate**: what to save to long-term memory
 - **Output gate**: what to return at the current time step
- Outputs h and v are the same


LSTM in PyTorch

```
class Net(nn.Module):
    def __init__(self, input_size):
        super().__init__()
        self.lstm = nn.LSTM(
            input_size=1,
            hidden_size=32,
            num_layers=2,
            batch_first=True,
        )
        self.fc = nn.Linear(32, 1)

    def forward(self, x):
        h0 = torch.zeros(2, x.size(0), 32)
        c0 = torch.zeros(2, x.size(0), 32)
        out, _ = self.lstm(x, (h0, c0))
        out = self.fc(out[:, -1, :])
        return out
```

- `__init__()` :
 - Replace `nn.RNN` with `nn.LSTM`
- `forward()` :
 - Add another hidden state `c`
 - Initialize `c` and `h` with zeros
 - Pass both hidden states to `lstm` layer

GRU cell



- Simplified version of LSTM cell
- Just one hidden state
- No output gate

GRU in PyTorch


```
class Net(nn.Module):
    def __init__(self, input_size):
        super().__init__()
        self.gru = nn.GRU(
            input_size=1,
            hidden_size=32,
            num_layers=2,
            batch_first=True,
        )
        self.fc = nn.Linear(32, 1)

    def forward(self, x):
        h0 = torch.zeros(2, x.size(0), 32)
        out, _ = self.gru(x, h0)
        out = self.fc(out[:, -1, :])
        return out
```

- `__init__()` :
 - Replace `nn.RNN` with `nn.GRU`
- `forward()` :
 - Use the `gru` layer

Should I use RNN, LSTM, or GRU?

- RNN is not used much anymore
- GRU is simpler than LSTM = less computation
- Relative performance varies per use-case
- Try both and compare



Let's practice!

INTERMEDIATE DEEP LEARNING WITH PYTORCH

Training and evaluating RNNs

INTERMEDIATE DEEP LEARNING WITH PYTORCH



Michał Oleszak
Machine Learning Engineer

Mean Squared Error Loss

- Error:

$$prediction - target$$

- Squared Error:

$$(prediction - target)^2$$

- Mean Squared Error:

$$\text{avg}[(prediction - target)^2]$$

Squaring the error:

- Ensures positive and negative errors don't cancel out
- Penalizes large errors more
- In PyTorch:

```
criterion = nn.MSELoss()
```

Expanding tensors

- Recurrent layers expect input shape
(batch_size, seq_length, num_features)
- We got (batch_size, seq_length)
- We must add one dimension at the end

```
for seqs, labels in dataloader_train:  
    print(seqs.shape)
```

```
torch.Size([32, 96])
```

```
seqs = seqs.view(32, 96, 1)  
print(seqs.shape)
```

```
torch.Size([32, 96, 1])
```

Squeezing tensors

- In evaluation loop, we need to revert the reshaping done in the training loop
- Labels are of shape (batch_size)

```
for seqs, labels in test_loader:  
    print(labels.shape)
```

torch.Size([32])
- Model outputs are (batch_size, 1)

```
out = net(seqs)
```

torch.Size([32, 1])
- Shapes of model outputs and labels must match for the loss function
- We can drop the last dimension from model outputs

```
out = net(seqs).squeeze()
```

torch.Size([32])

Training loop

```
net = Net()  
criterion = nn.MSELoss()  
optimizer = optim.Adam(  
    net.parameters(), lr=0.001  
)  
  
for epoch in range(num_epochs):  
    for seqs, labels in dataloader_train:  
        seqs = seqs.view(32, 96, 1)  
        outputs = net(seqs)  
        loss = criterion(outputs, labels)  
        optimizer.zero_grad()  
        loss.backward()  
        optimizer.step()
```

- Instantiate model, define loss & optimizer
- Iterate over epochs and data batches
- Reshape input sequence
- The rest: as usual

Evaluation loop

```
mse = torchmetrics.MeanSquaredError()

net.eval()
with torch.no_grad():
    for seqs, labels in test_loader:
        seqs = seqs.view(32, 96, 1)
        outputs = net(seqs).squeeze()
        mse(outputs, labels)

print(f"Test MSE: {mse.compute()}")
```

Test MSE: 0.13292162120342255

- Set up MSE metric
- Iterate through test data with no gradients
- Reshape model inputs
- Squeeze model outputs
- Update the metric
- Compute final metric value

LSTM vs. GRU

- LSTM:

```
Test MSE: 0.13292162120342255
```

- GRU:

```
Test MSE: 0.12187089771032333
```

- GRU preferred: same or better results with less processing power

Let's practice!

INTERMEDIATE DEEP LEARNING WITH PYTORCH